

ForestClaw : Ghost filling and parallel communication

Donna Calhoun (Boise State University)

Carsten Burstedde, Univ. of Bonn, Germany

p4est Summer School

July 20 - 25, 2020

Bonn, Germany (Virtual)

p4est interface for ForestClaw

Files that provide the interface between p4est and ForestClaw

- **src/forestclaw2d.h** - definitions of *patch*, *block*, and *domain* structs.
- **src/forestclaw2d.c** - nearest neighbors searches, transformations for multi-block boundaries, iterators, tagging
- **src/fclaw_base.c** - option handling utilities
- **src/fclaw2d_convenience.c** - multi-block domain definitions, routines for adapting and partitioning the domain
- Additional header files, and a few more files that provide mapping utilities for cubed sphere, torus, and so on.

ForestClaw is built on top of the routines in these files.

Face neighbor searches

```
fclaw2d_patch_relation_t  
  fclaw2d_patch_face_neighbors(fclaw2d_domain_t * domain,  
                                int blockno, int patchno, int faceno,  
                                int rproc[P4EST_HALF], int *rblockno,  
                                int rpatchno[P4EST_HALF], int *rfaceno)  
{  
    /* Returns neighbor type (BOUNDARY, HALFSIZE, SAMESIZE, DOUBLESIZE) */  
  
    /* Additional output : MPI rank, patch number and block number for  
       remote patch neighbors. */  
}
```

- This is one of two essential routines needed to build ghost-filling infrastructure for ForestClaw.

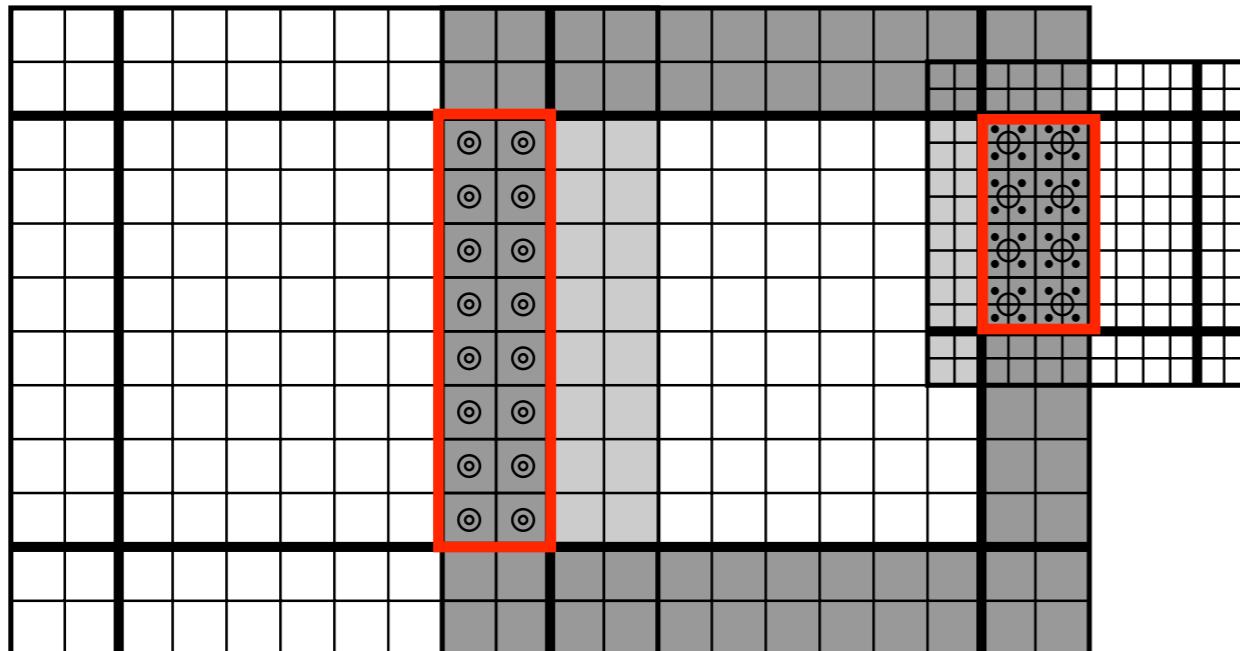
Corner neighbor searches

```
int  
fclaw2d_patch_corner_neighbors (fclaw2d_domain_t * domain,  
                                int blockno, int patchno, int cornerno,  
                                int *rproc, int *rblockno, int *rpatchno,  
                                int *rcorner,  
                                fclaw2d_patch_relation_t * neighbor_size)  
{  
    /* Returns 0,1 to indicate whether patch has a corner neighbor. */  
}
```

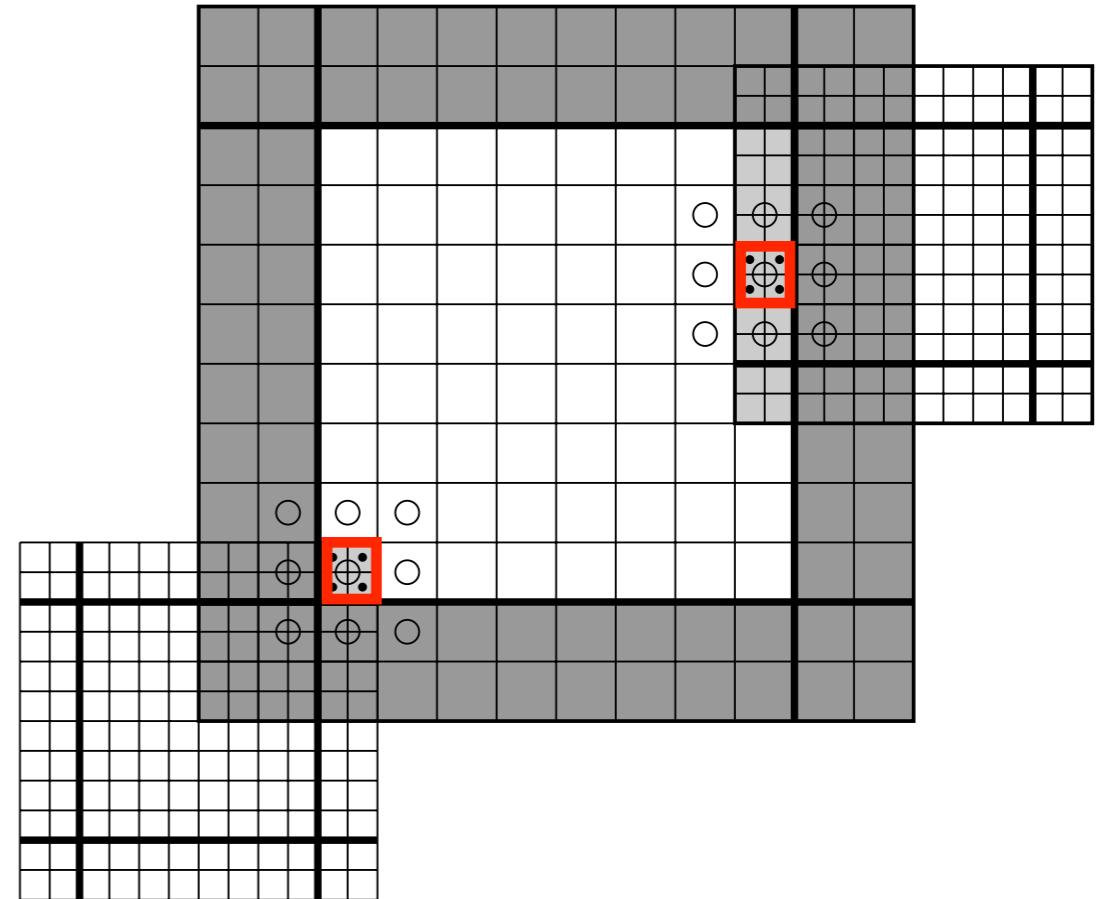
- Corner information needed for unsplit finite volume schemes.
- Corners exchange introduced some new challenges for parallel communication in p4est.

Filling ghost cells

Assume valid data in the interior of each patch



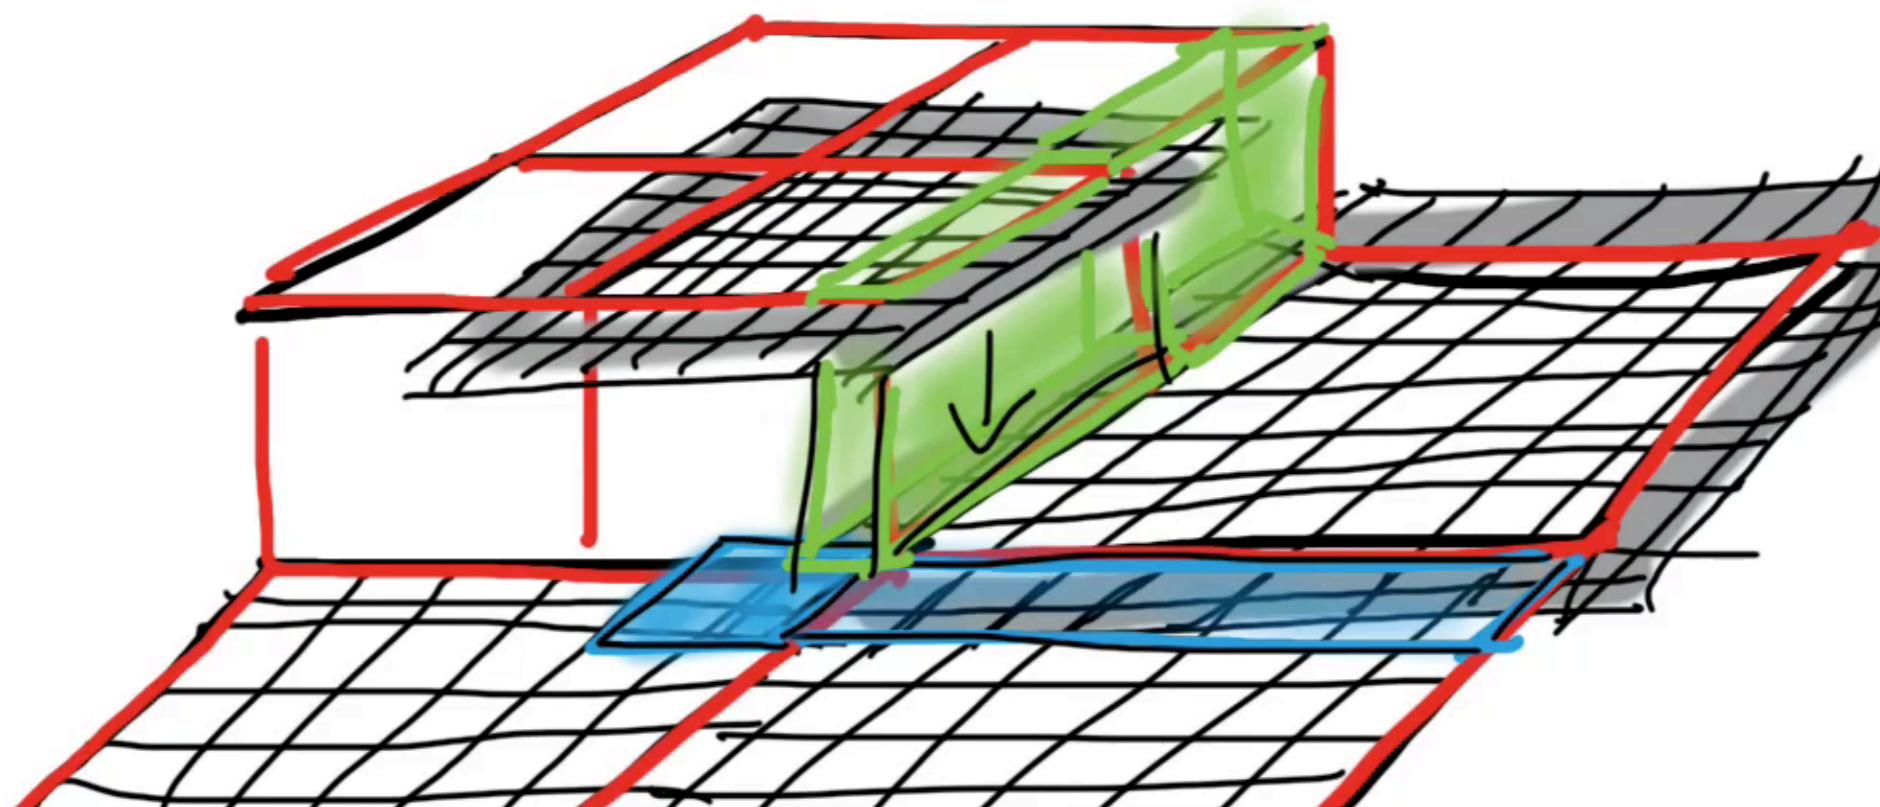
Step 1 : Fill “coarse grid” cells.
Copy between same size
neighbors; Average from fine grid
to coarse grid.



Step 2 : Interpolate from coarse
grid to fine ghost regions, using
coarse grid ghost regions

*Unsplit version of finite volume wave-propagation algorithm requires corner exchanges.
Two layers of ghost cells needed for limiting waves to avoid unphysical oscillations.*

How are ghost cells filled?

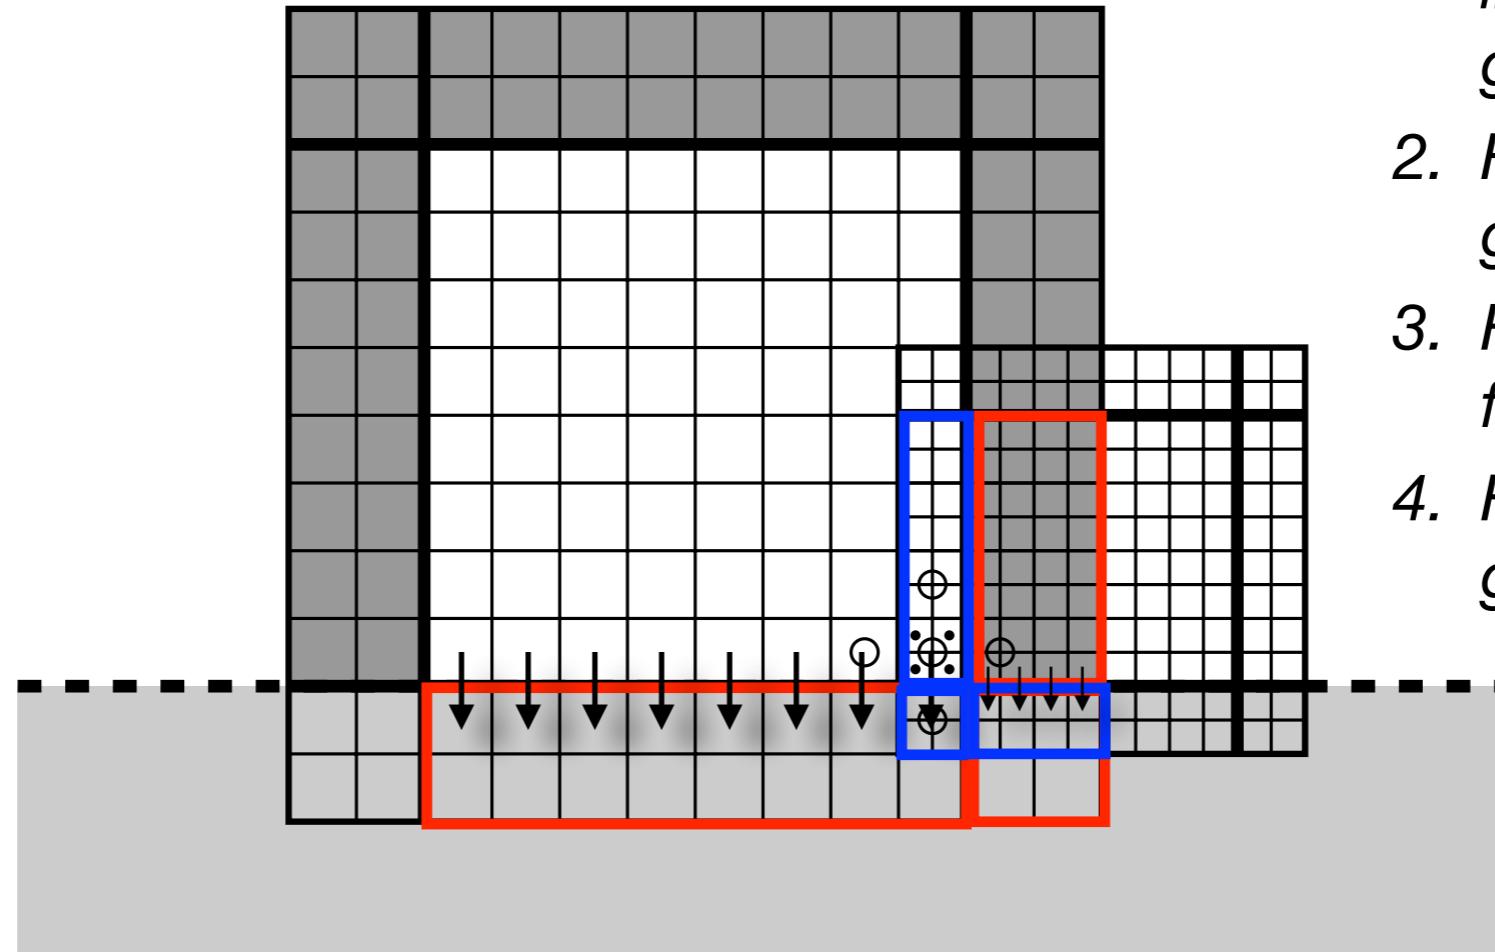


Copy

Average

Interpolate

Physical boundary conditions



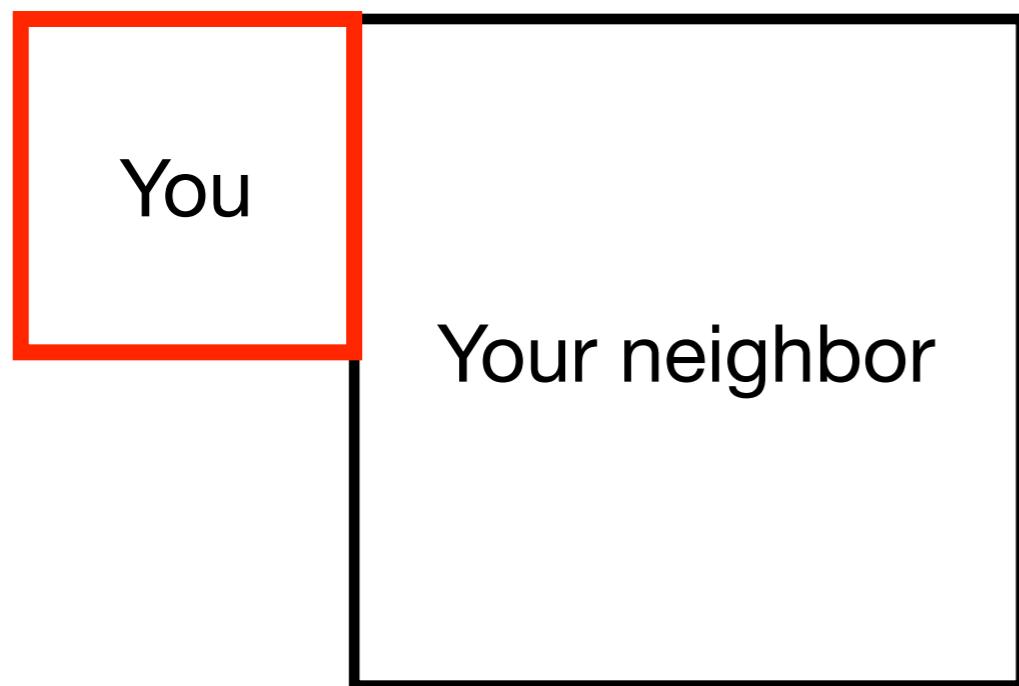
1. Fill exterior and interior coarse face ghost regions
2. Fill exterior coarse grid corner region
3. Fill fine grid interior face region
4. Fill exterior corner ghost regions

- Two passes of physical boundary conditions are required to fill corners in the exterior region.

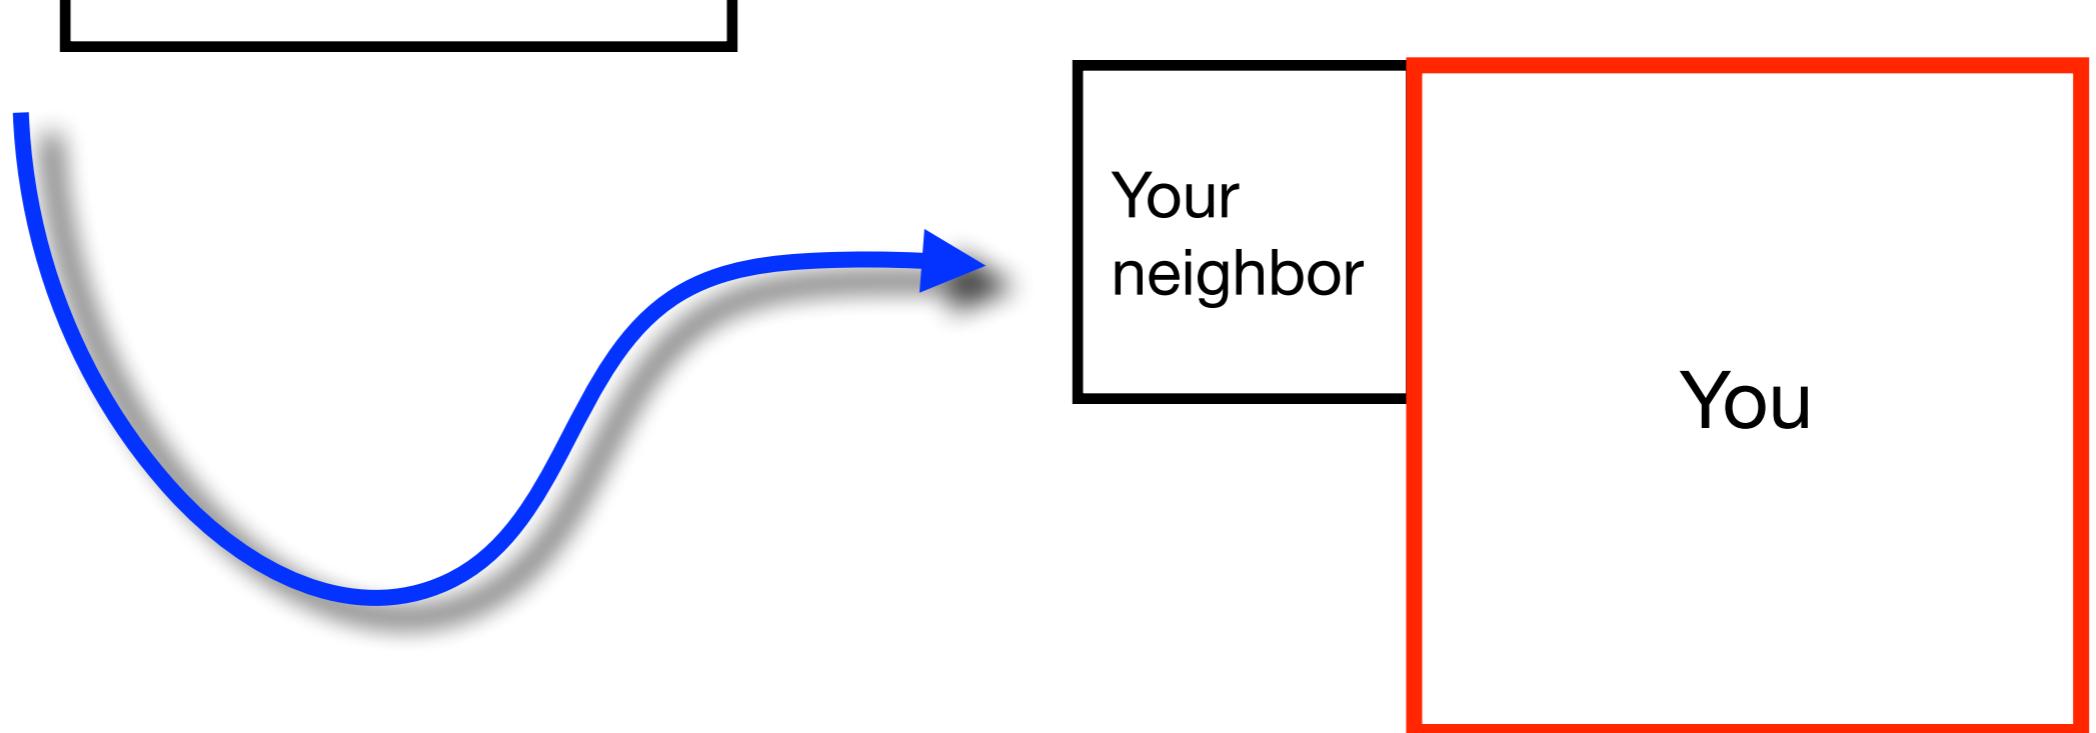
Ghost filling

- Iterators used to iterate over patches
 - Sequencing very important, so multiple iterations over ghost cells are required.
- For each patch, nearest neighbors are queried.
 - Depending on stage in sequence, face or corner ghost regions may or may not be filled.
- 20 possible arrangements of a grid and neighbors (not including potential rotations at multi block boundaries) reduced to 12. Trick : A grid with a double size neighbor is swapped with its neighbor.
- Routines for ghost-filling at faces are parameterized by direction (0,1), face (0,1,2,3), and neighbor type, so that only three routines are needed - one for copying, one for averaging, and one for interpolation.
- Routines for ghost-filling at corners are parameterized by corner number and neighbor type. Three routines for copying, averaging and interpolation

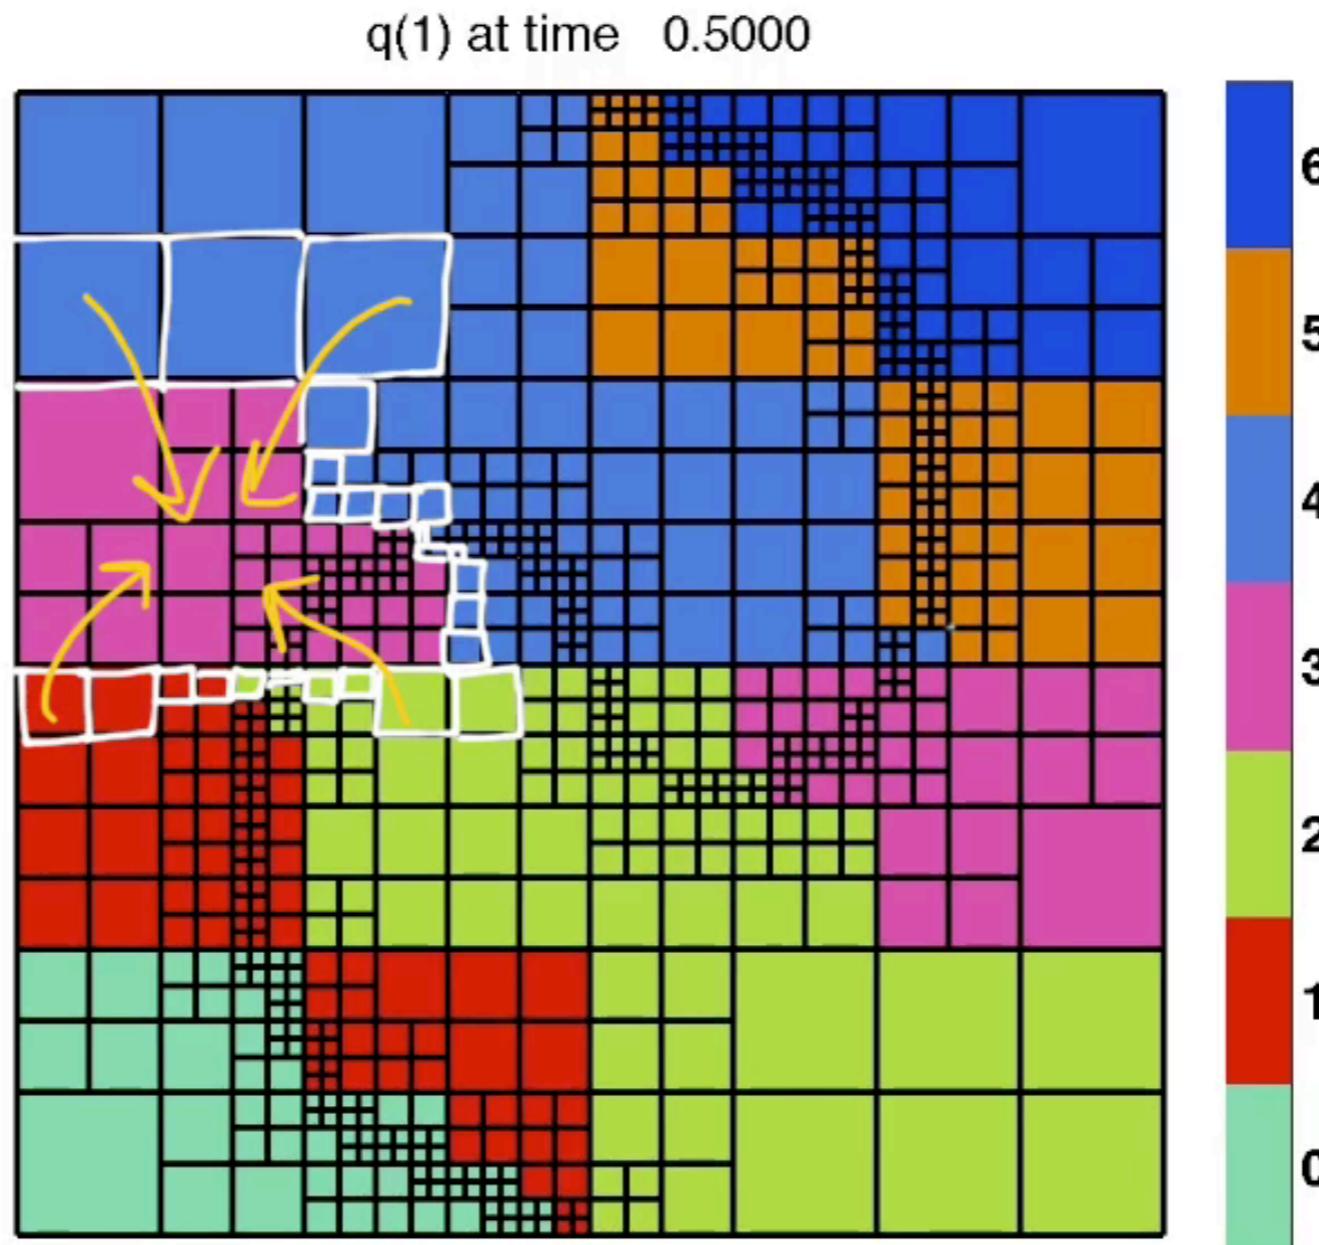
“Context switching”



- “Context switching” allows us to reduce possible combinations of grid pairings.
- Uses a “swap” routine supplied by p4est so that face numbers are relative to “You” and not your neighbor.
- Works seamlessly with multi-block boundaries.

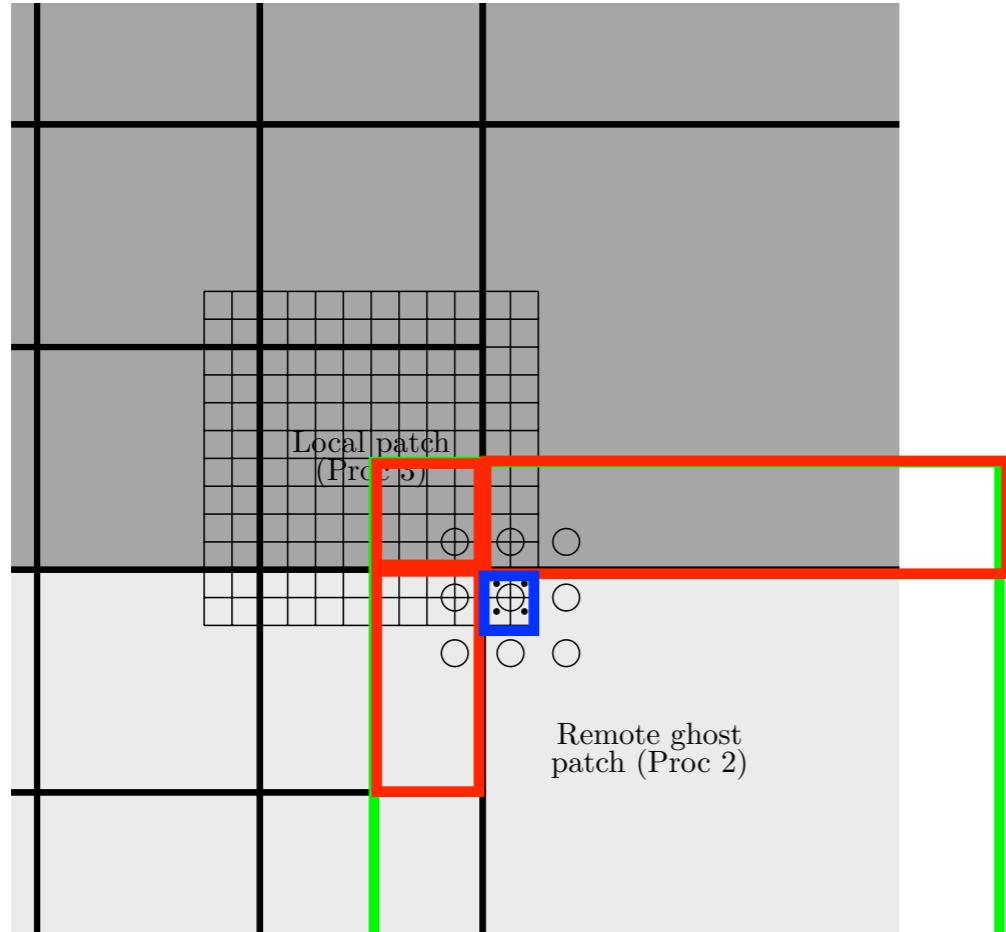


Parallel ghost filling

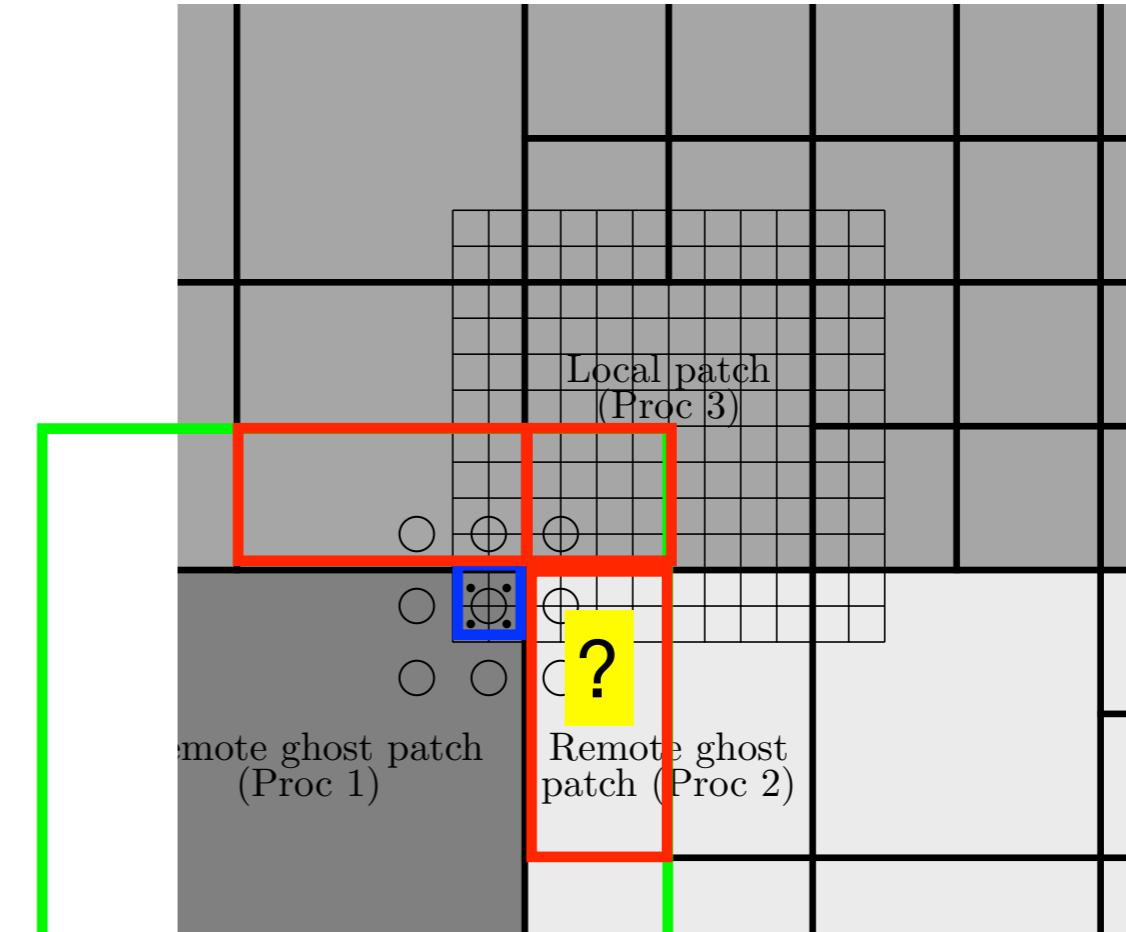


- Remote patches are created by p4est, and are stored in separate data structure
- Patch routines in ForestClaw are used to re-build essential information in ghost patches

Parallel ghost filling algorithm



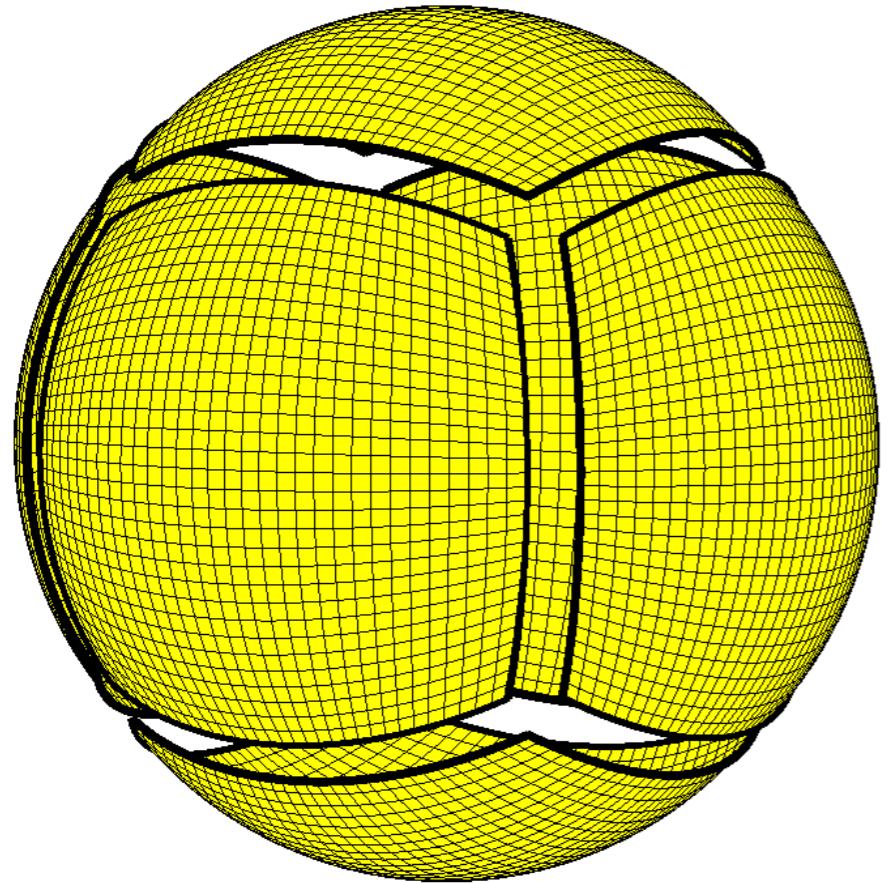
Remote patches on processor boundary must exchange ghost cells before being sent to local processor



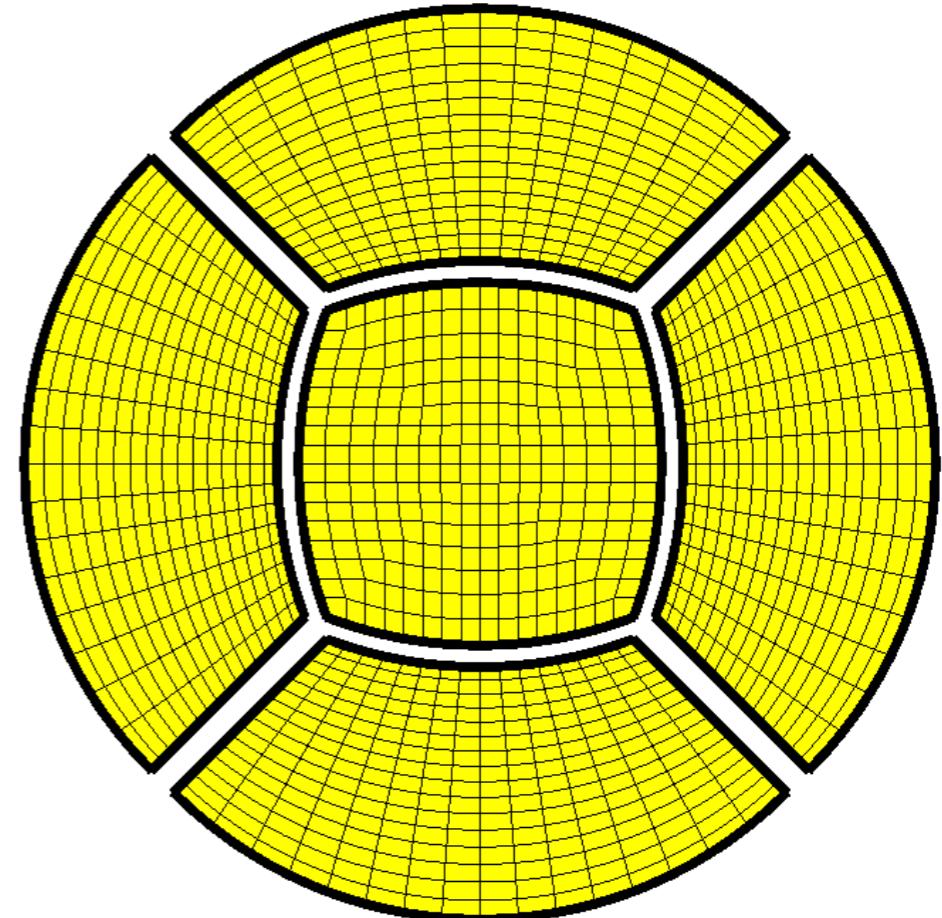
An lightweight indirect exchange is required between remote proc 2 and 3

- Remote patches must have valid coarse grid ghost data so that corners on local patches can be filled in.
- Requires one communication pass per ghost cell update

Multiblock boundaries



- Ghost filling at multi-block boundaries is transparent to the user
- Requires index transformations supplied by p4est interface
- Straightforward to modify coarse/fine averaging and interpolation stencils, even at multi-block boundaries.



Questions?

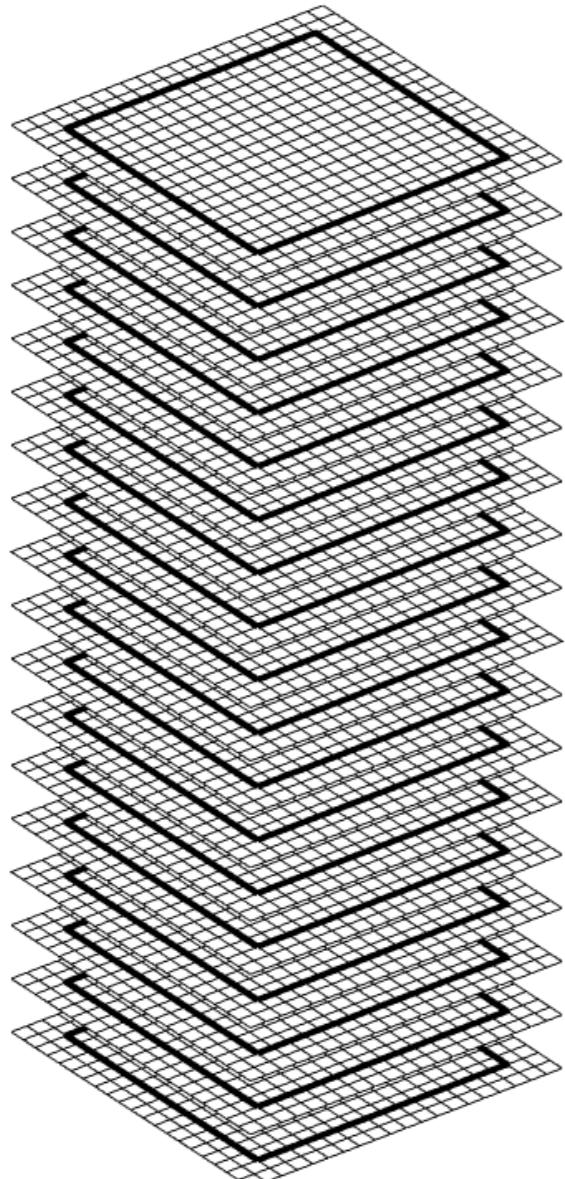
ForestClaw on GPUs

Ported fully unsplit wave propagation algorithm for hyperbolic conservation laws (implemented in Clawpack) to CUDA.

- Copy time level solution on all patches to single contiguous block of CPU memory
- Copy contiguous block of CPU memory to the GPU.
- Configure the GPU to assign one 1d thread block to each single ForestClaw patch
- Divide shared memory equally among thread blocks=patches
- All solution data resides in global memory; shared memory is only used for temporary data
- CUDA function pointers used to provide custom Riemann solvers.
- Best to use the 4.x (SOA) data layout
- All core ForestClaw routines, and p4est remain on the CPU. Only the patch update is ported to the GPU.

ForestClaw on GPUs

```
dim3 grid(1,1,batch_size);
```



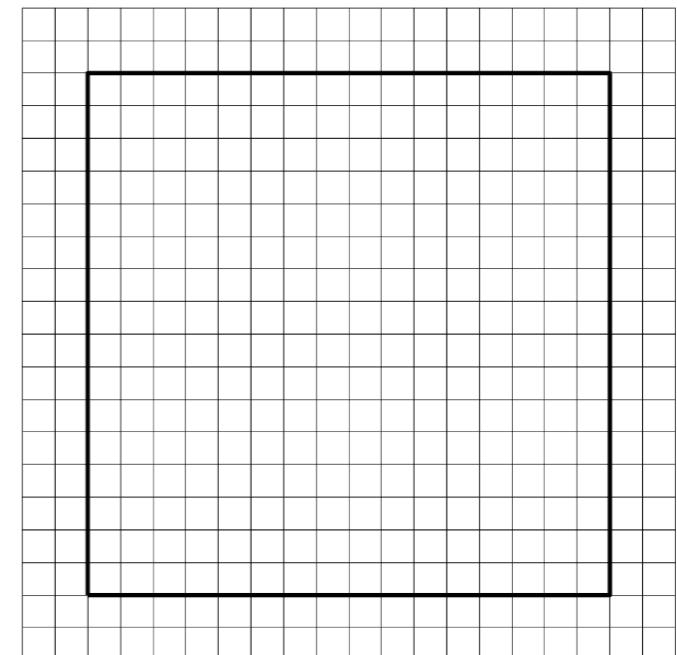
One ForestClaw patch per
CUDA block

```
block_size = 128; batch_size = 4000;  
mwork = 9*mEqn + 9*maux + mwaves + mEqn*mwaves;  
bytes_per_thread = sizeof(double)*mwork;  
bytes = bytes_per_thread*block_size;  
  
dim3 block(block_size,1,1);  
dim3 grid(1,1,batch_size);  
  
claw_flux2<<<grid,block,bytes>>>(mx,my,mEqn,...)
```

1d thread blocks
3d grid

~4000 patches in a batch
~128 threads per block

Patch layout with
valid ghost cell data



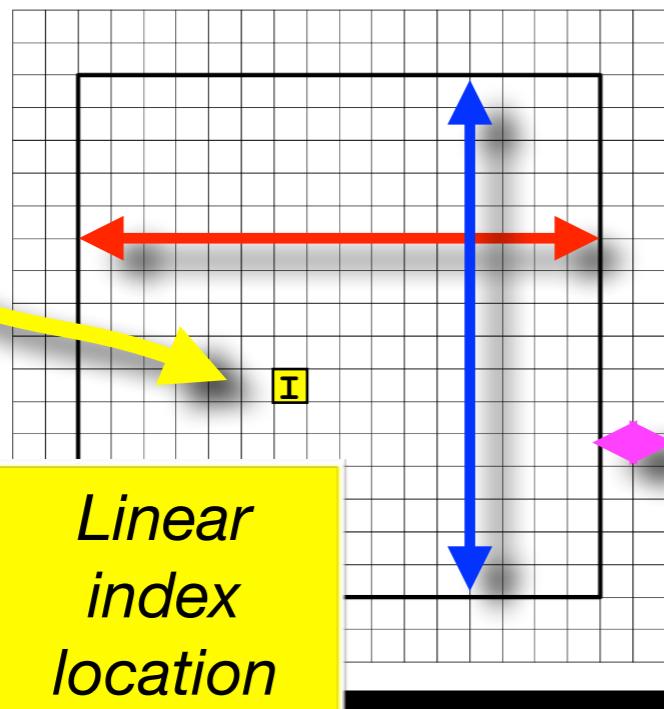
Thread block - loop over faces

```
ys = (2*mbc + mx); /* Stride */
ifaces_x = mx + 2*mbc-1;
ifaces_y = my + 2*mbc-1;
num_cells = ifaces_x*ifaces_y;

for(ti = threadIdx.x; ti < num_ifaces; ti += blockDim.x)
{
    ix = ti % ifaces_x;
    iy = ti/ifaces_x;

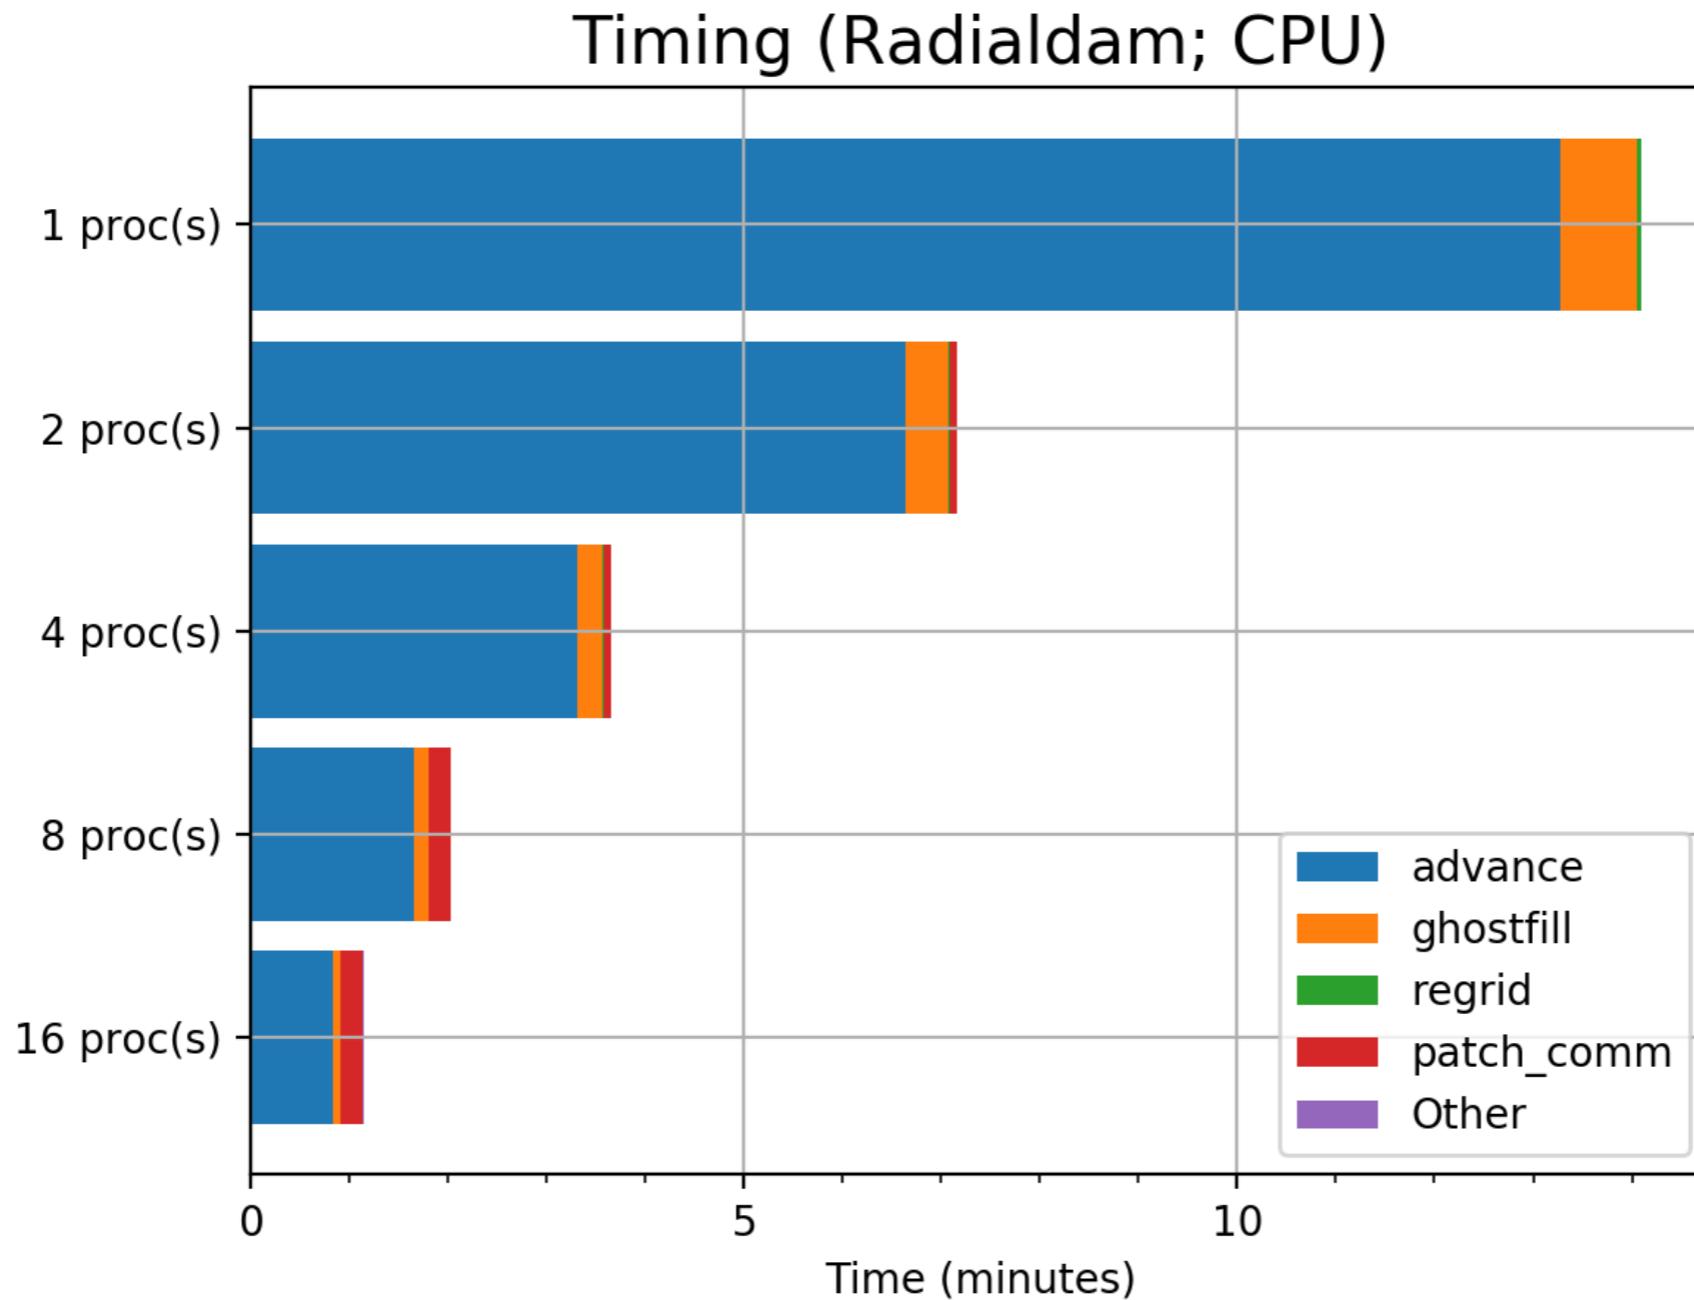
    I = (iy + 1)*ys + (ix + 1);
    ...
}
```

Solve a normal Riemann problem at each face;
include 1 ghost cell in each direction

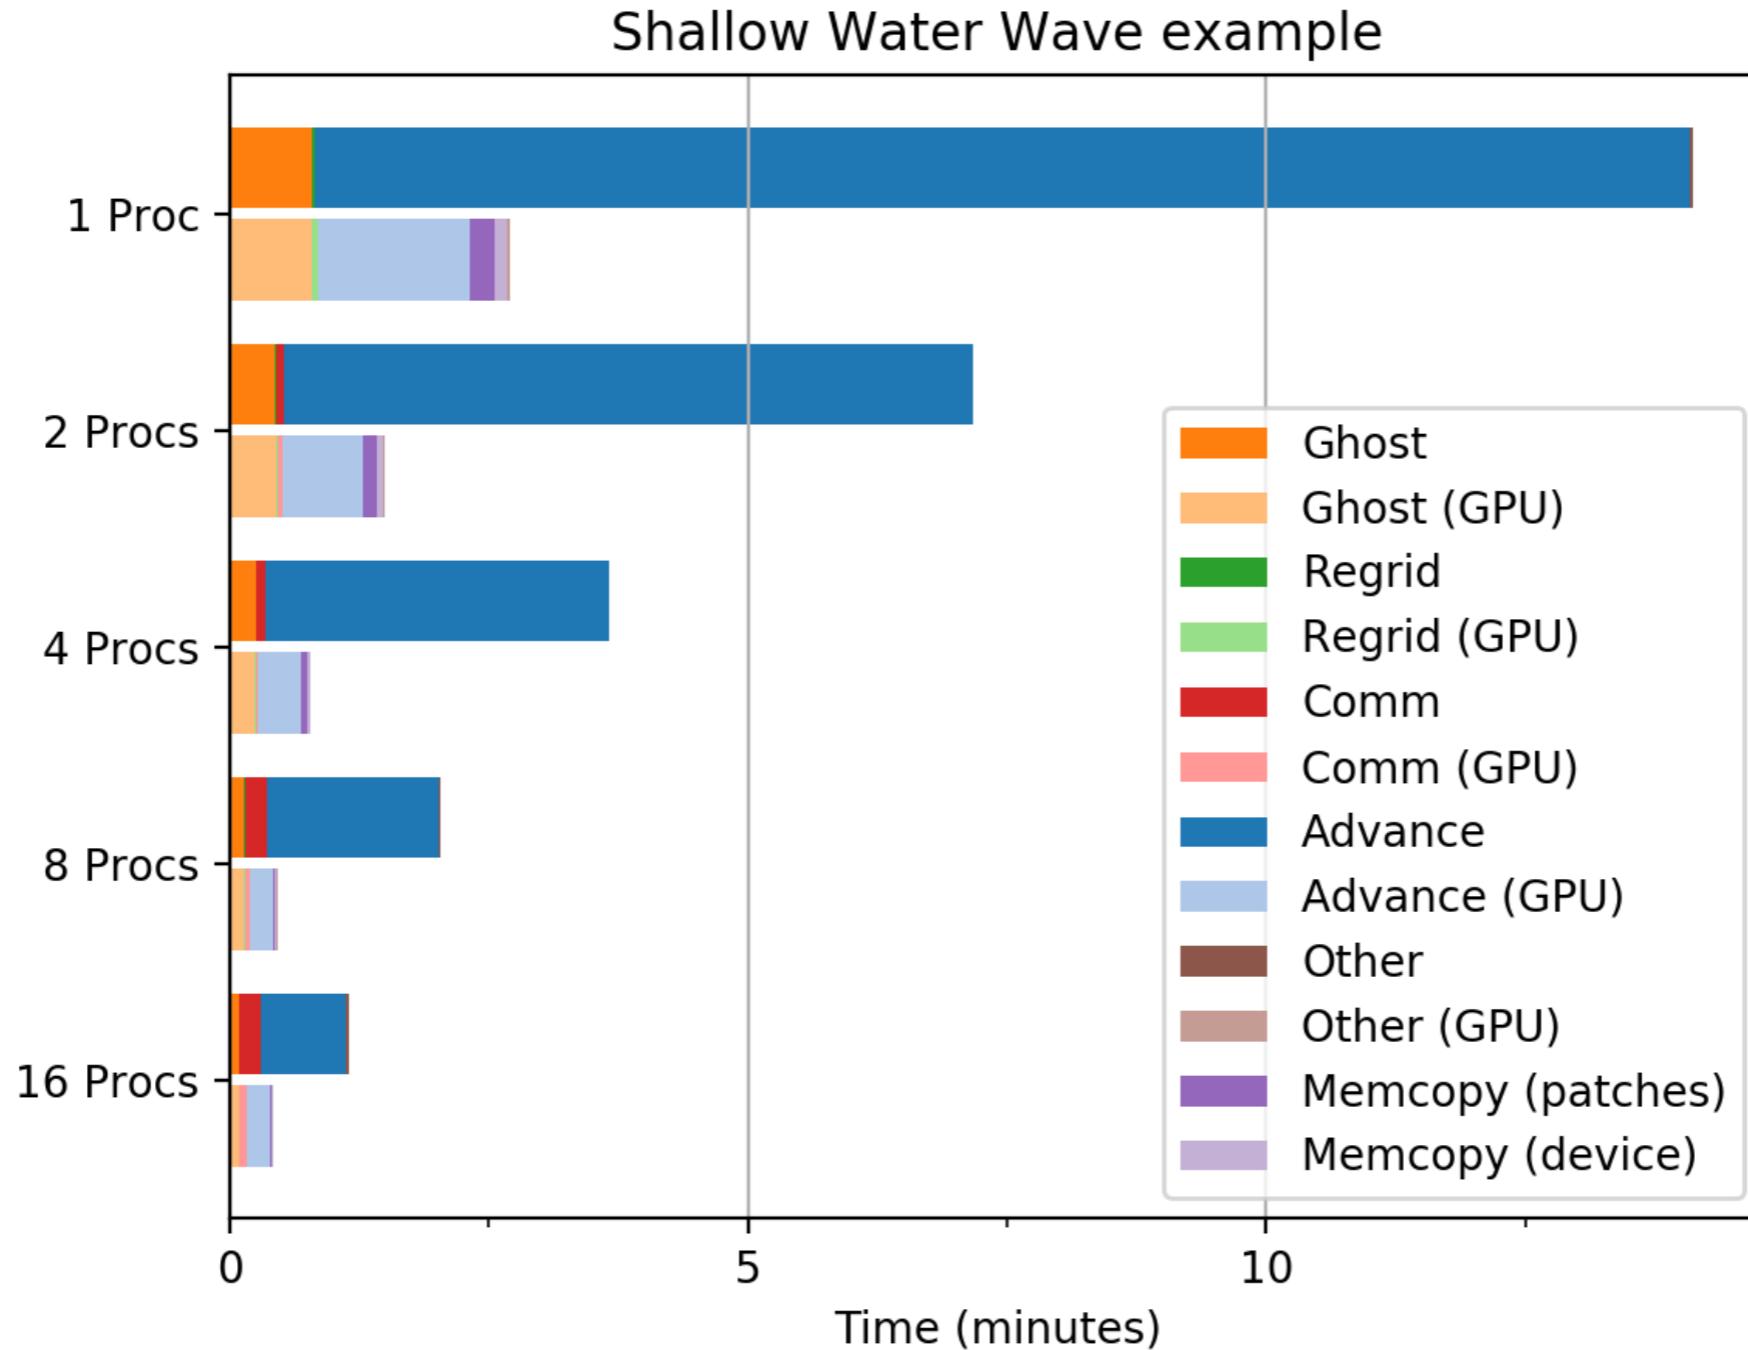


mx : Number of interior grid cells in x
my : Number of interior grid cells in y
mbc : Number of ghost cells

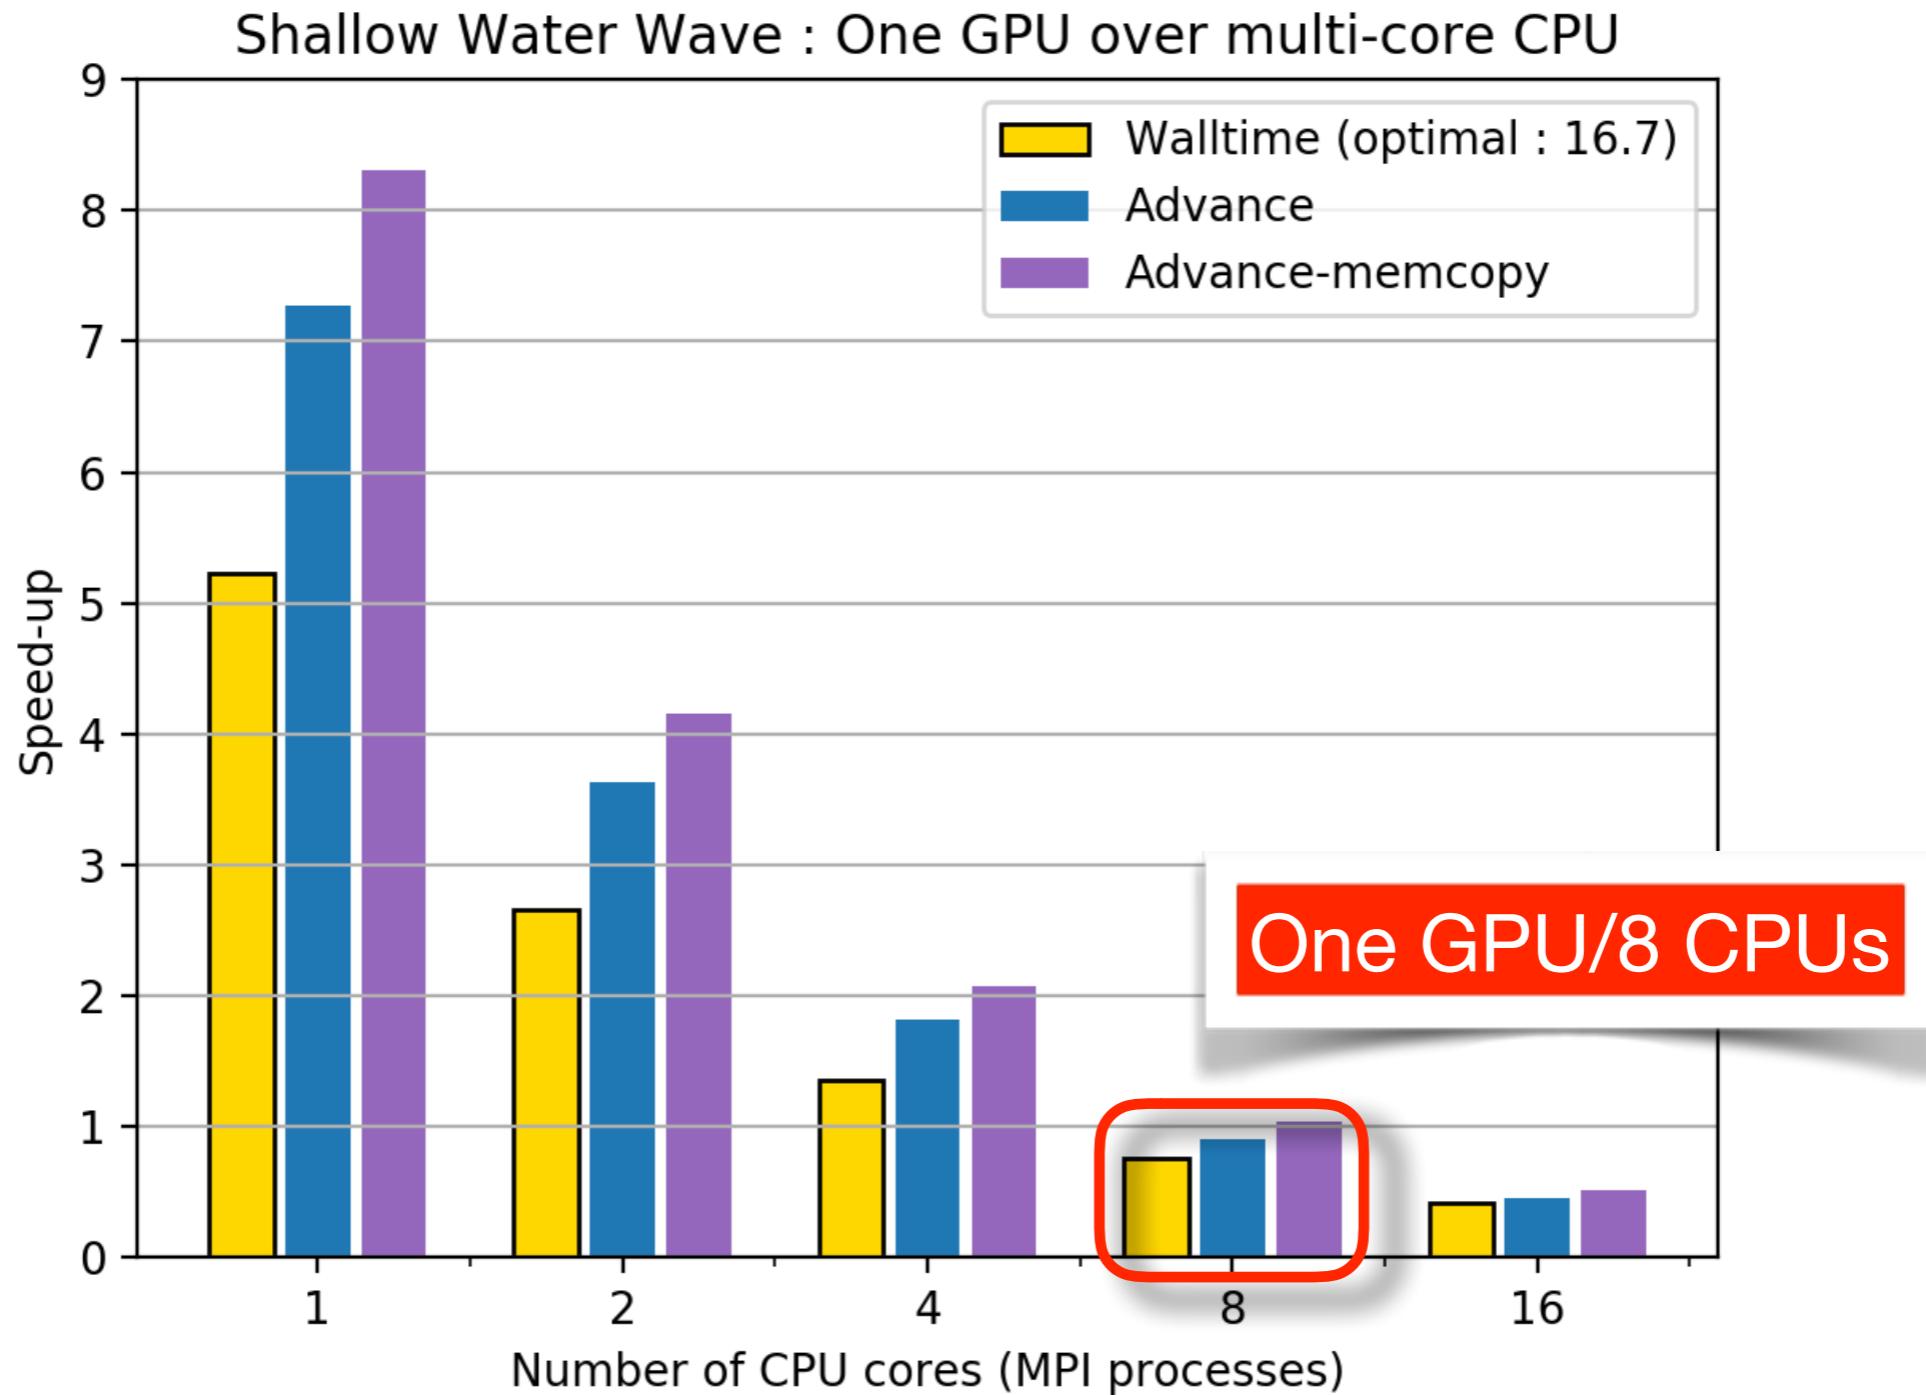
Shallow water



Shallow water

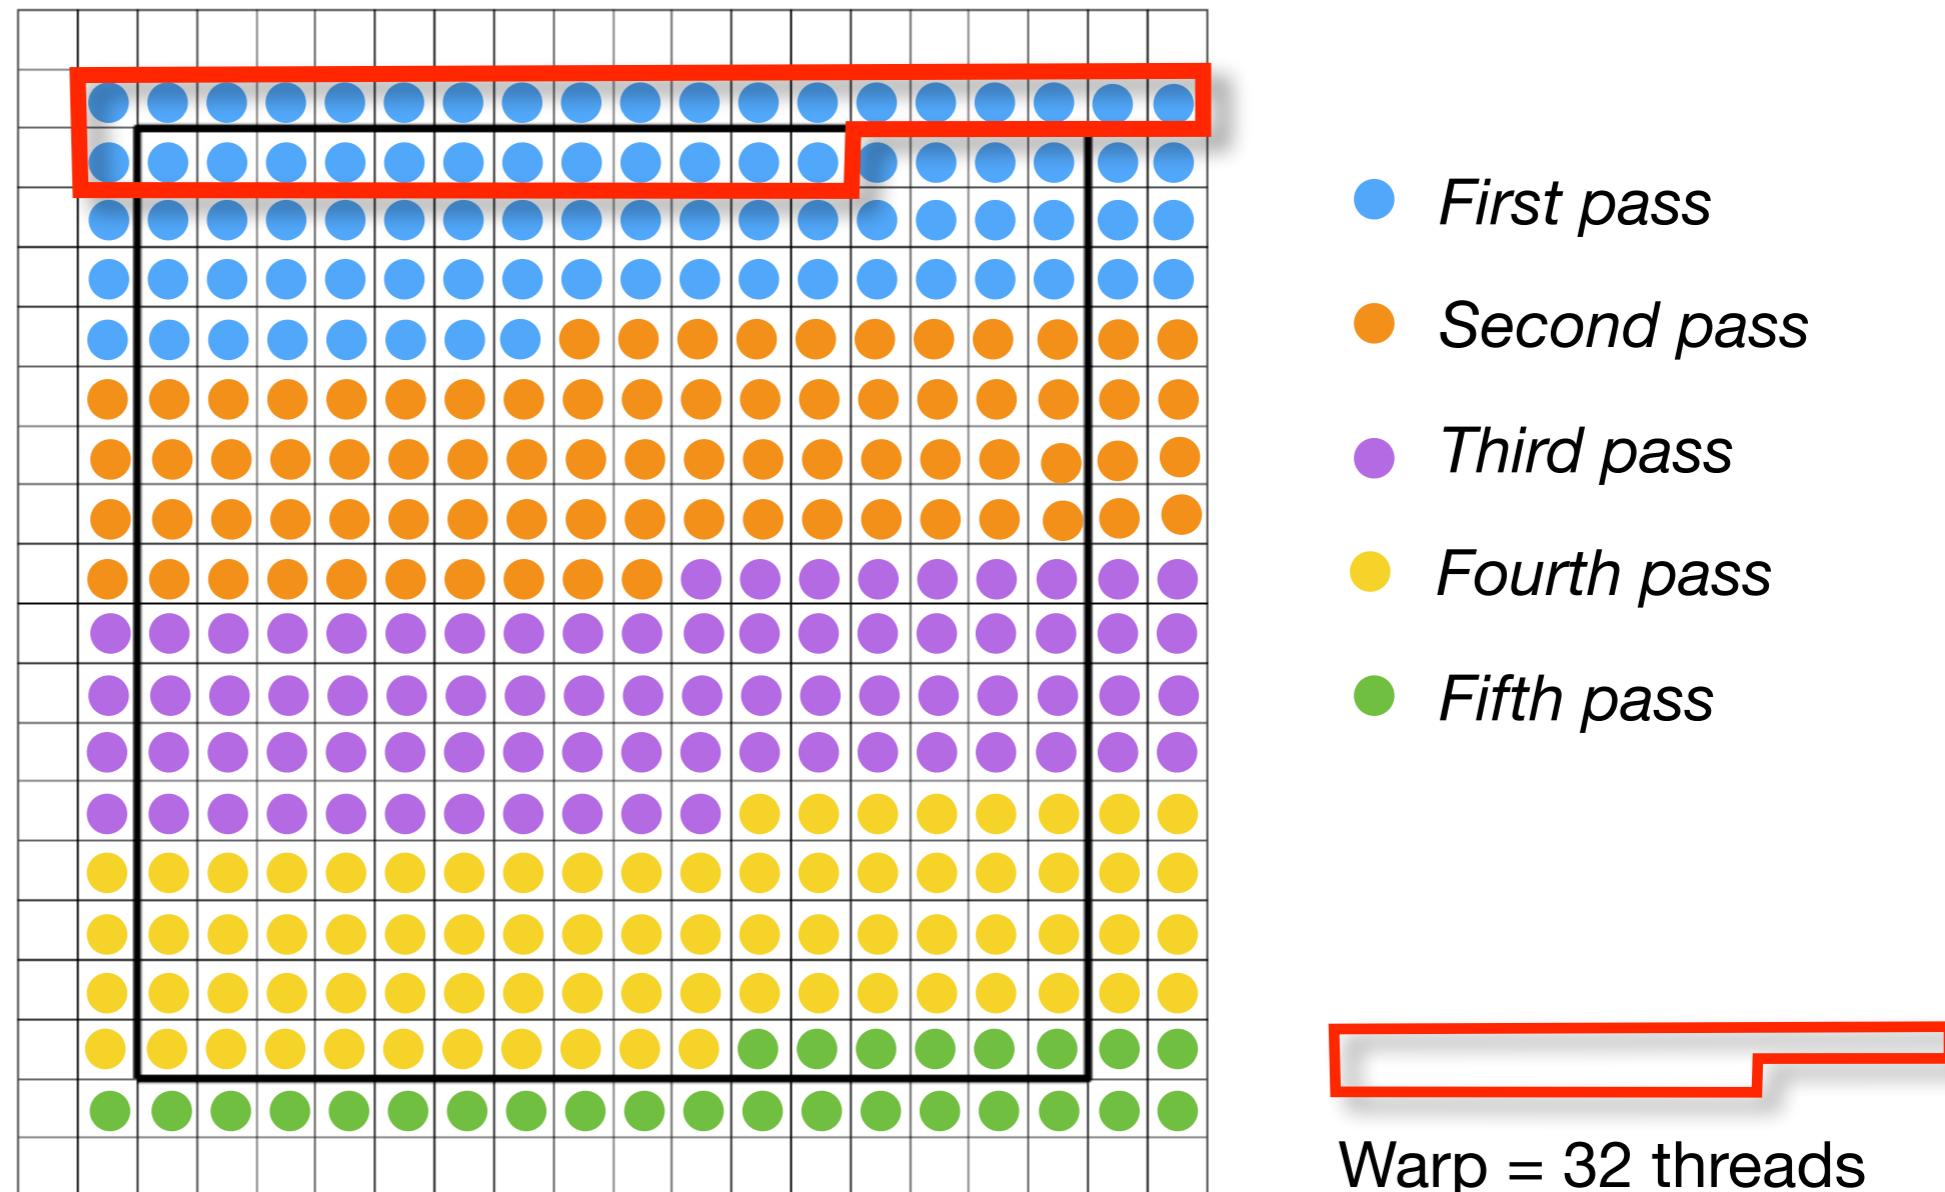


Shallow water



One dimensional thread block

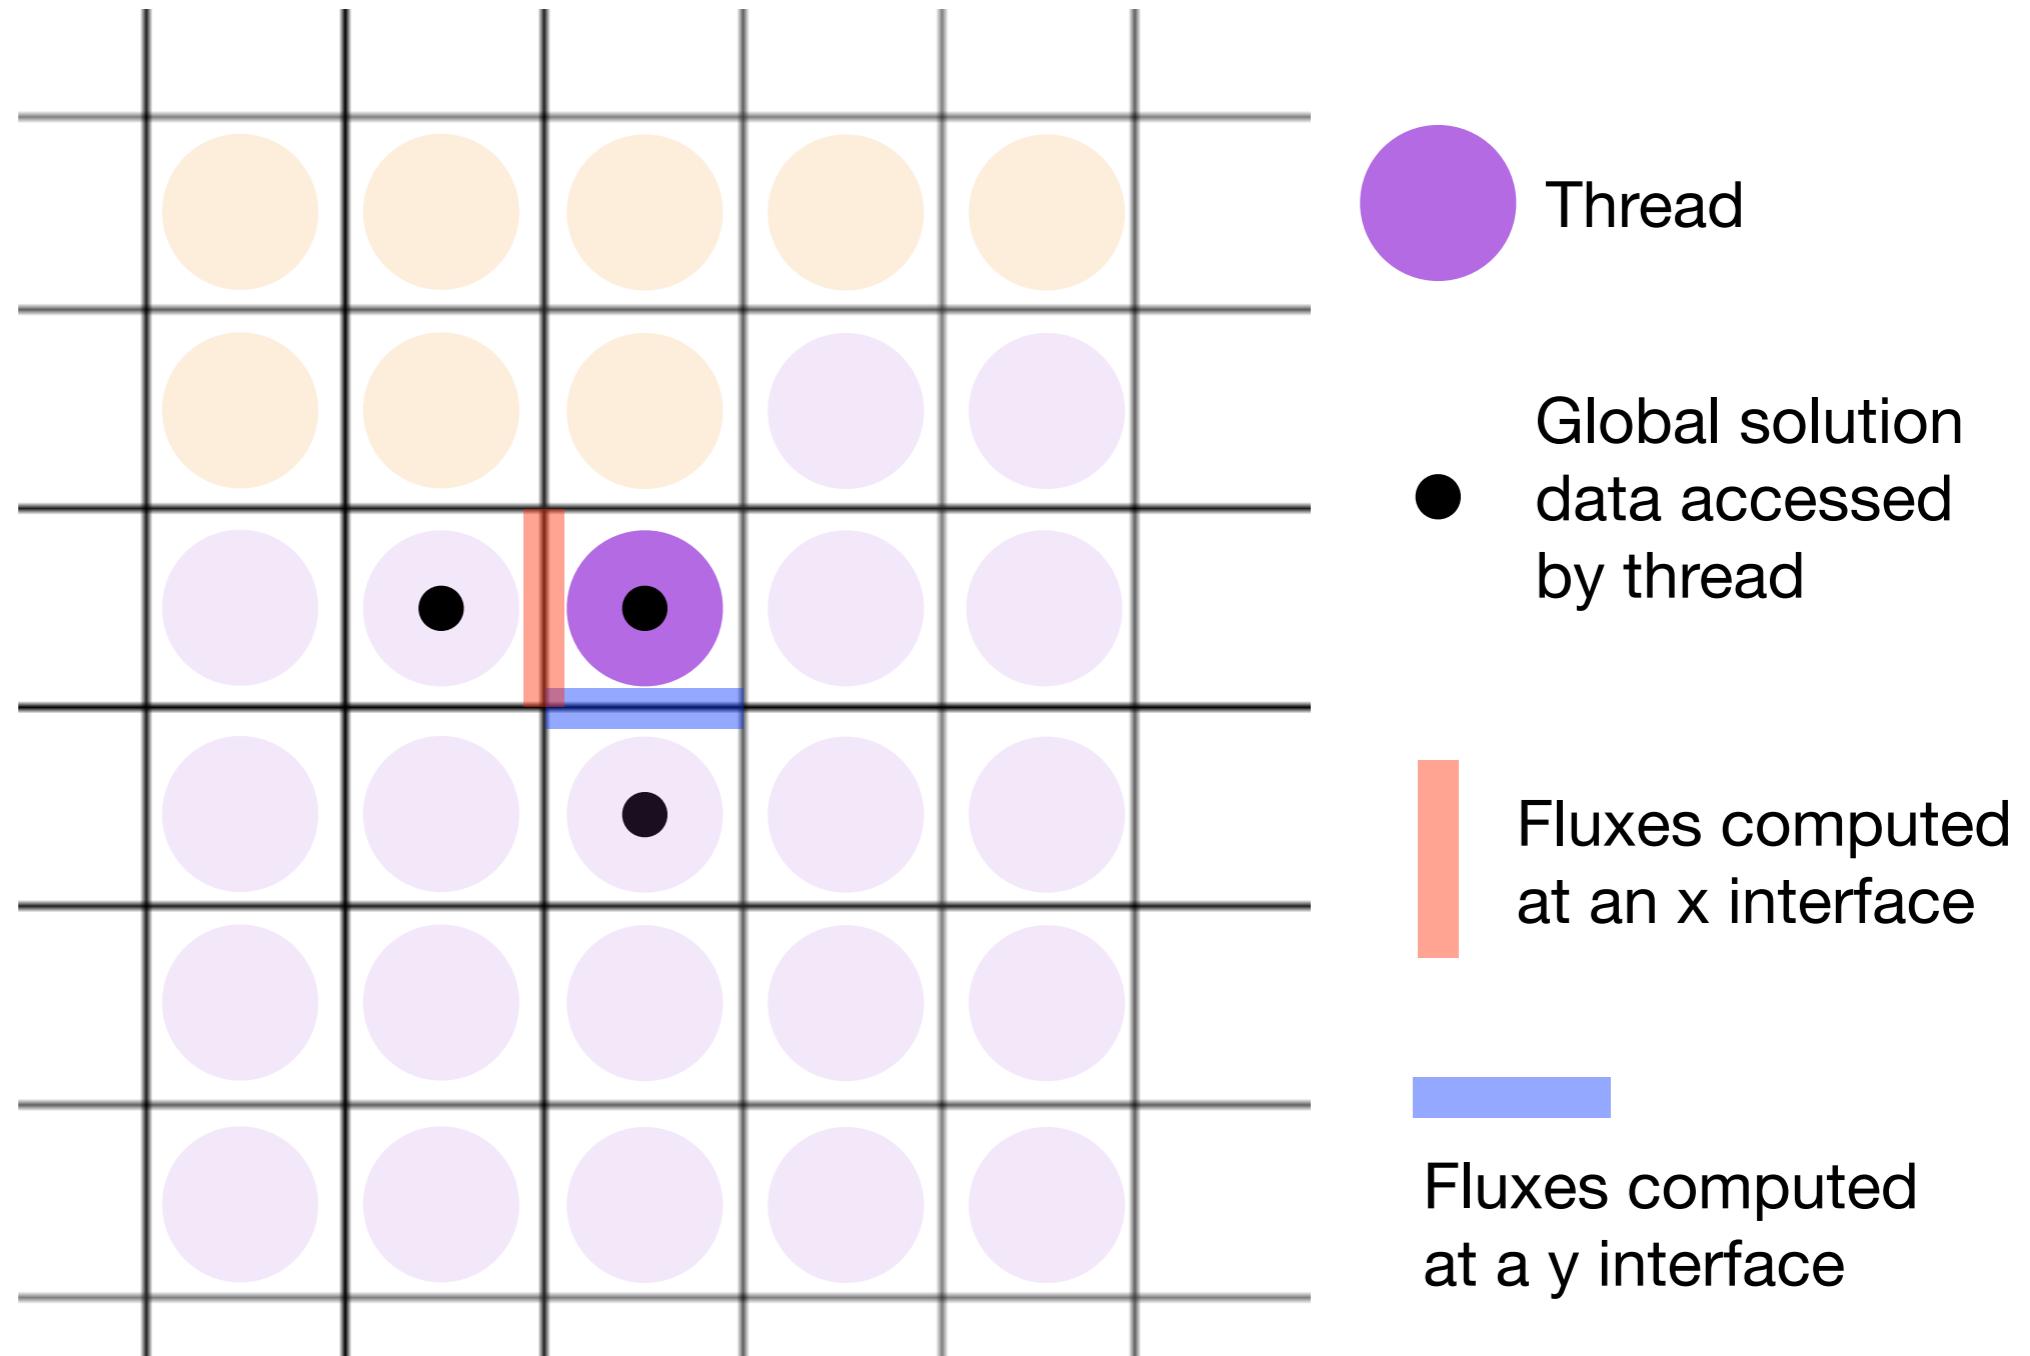
- No block synchronization required
- Typical patch sizes are 32x32
- Number of threads per patch : ~128, depending on shared memory requirements



Solve Riemann problems at x and y faces

Normal Riemann problems

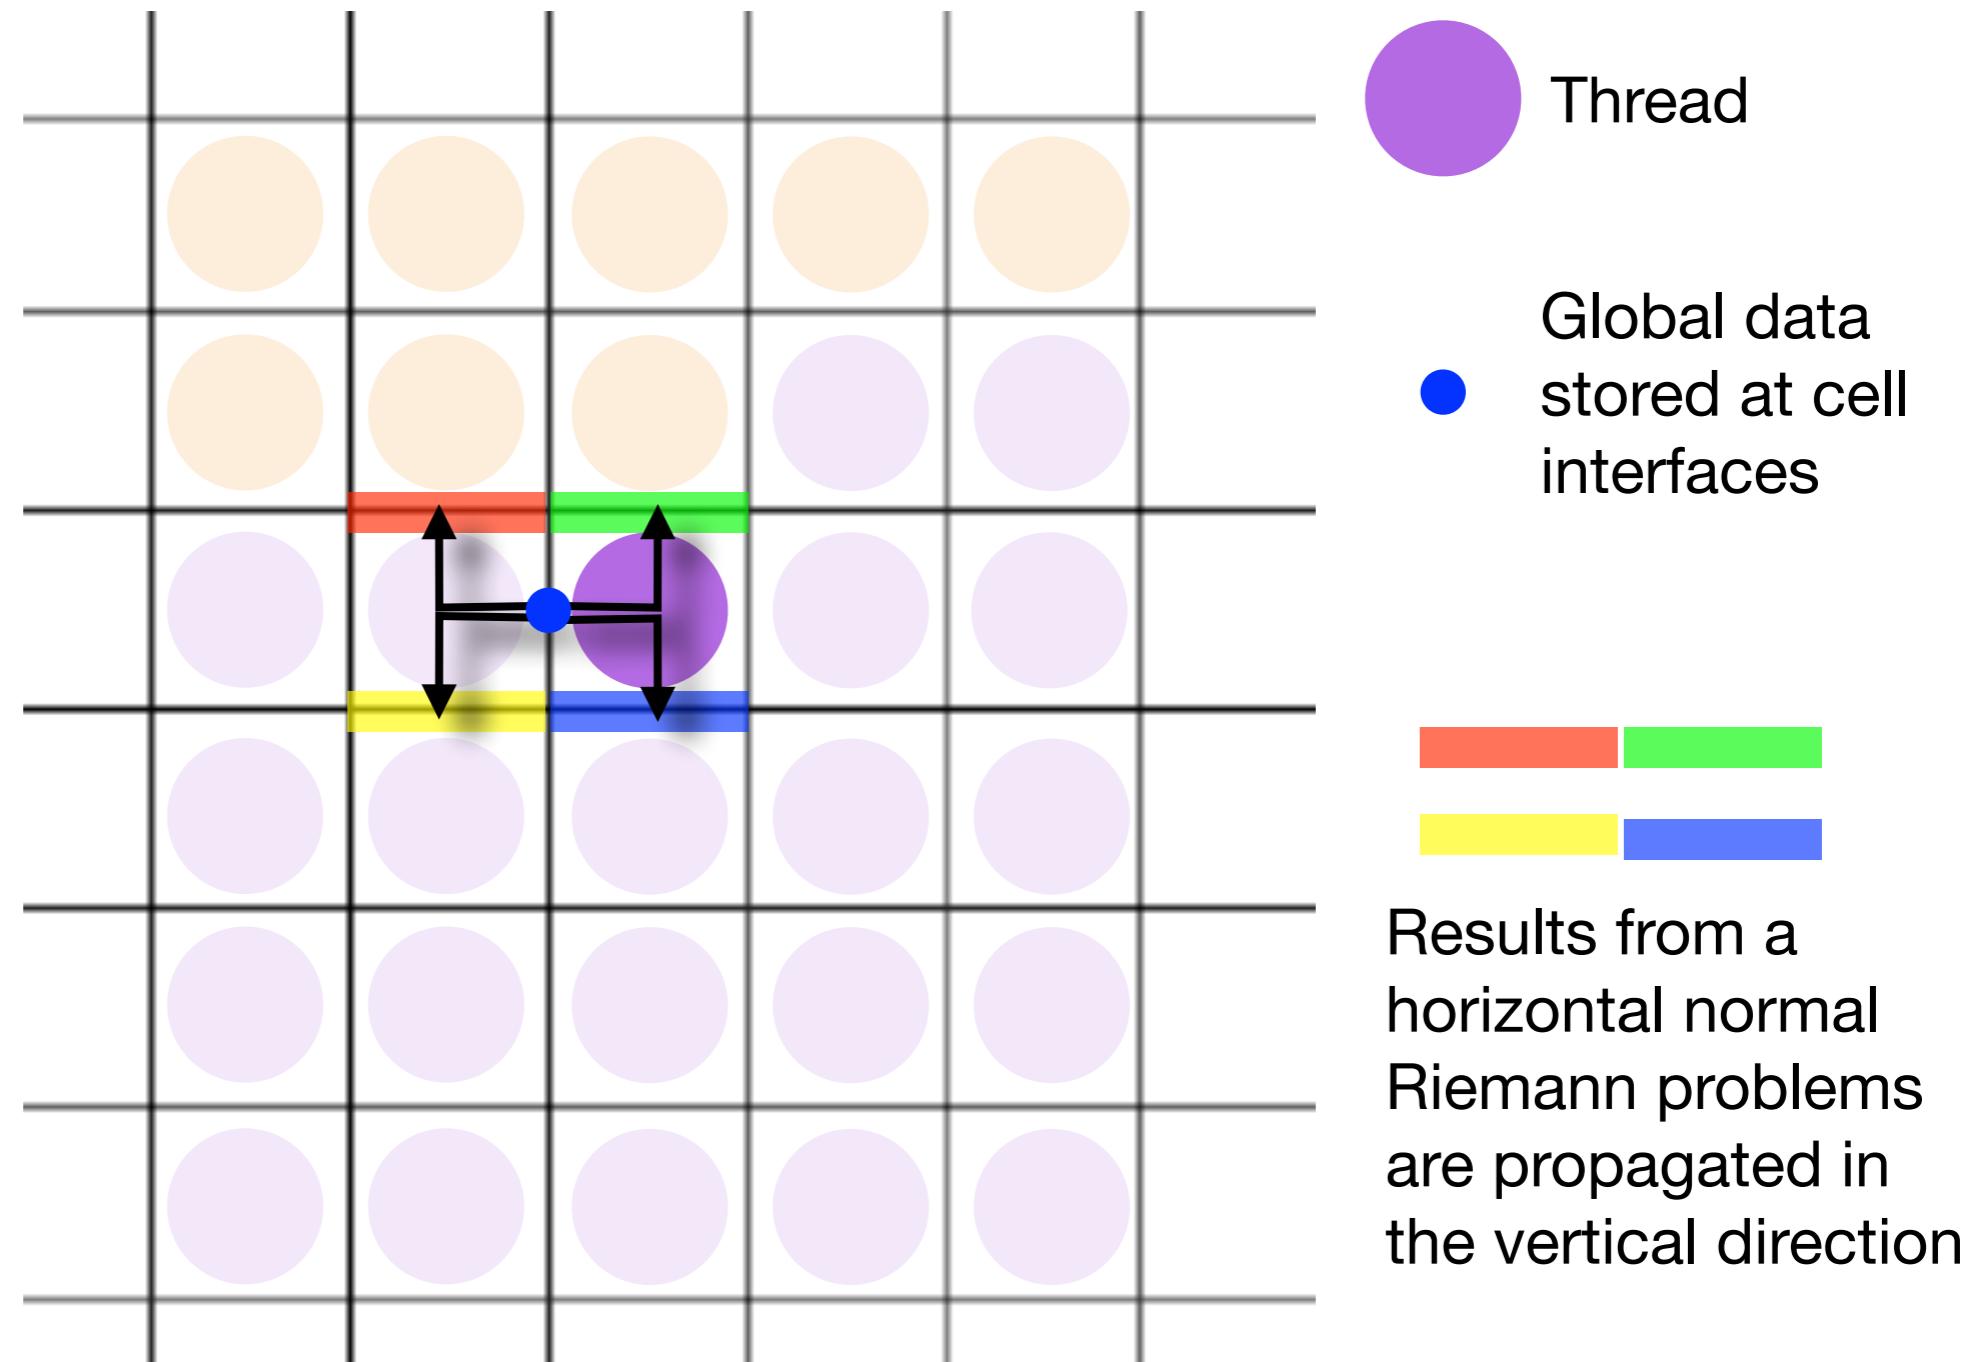
- Each thread makes a local copy of global data and stores it in shared memory.
- Fluxes computed Riemann problems stored in global array



Fluxes are computed by solving Riemann problems

Unsplit algorithm

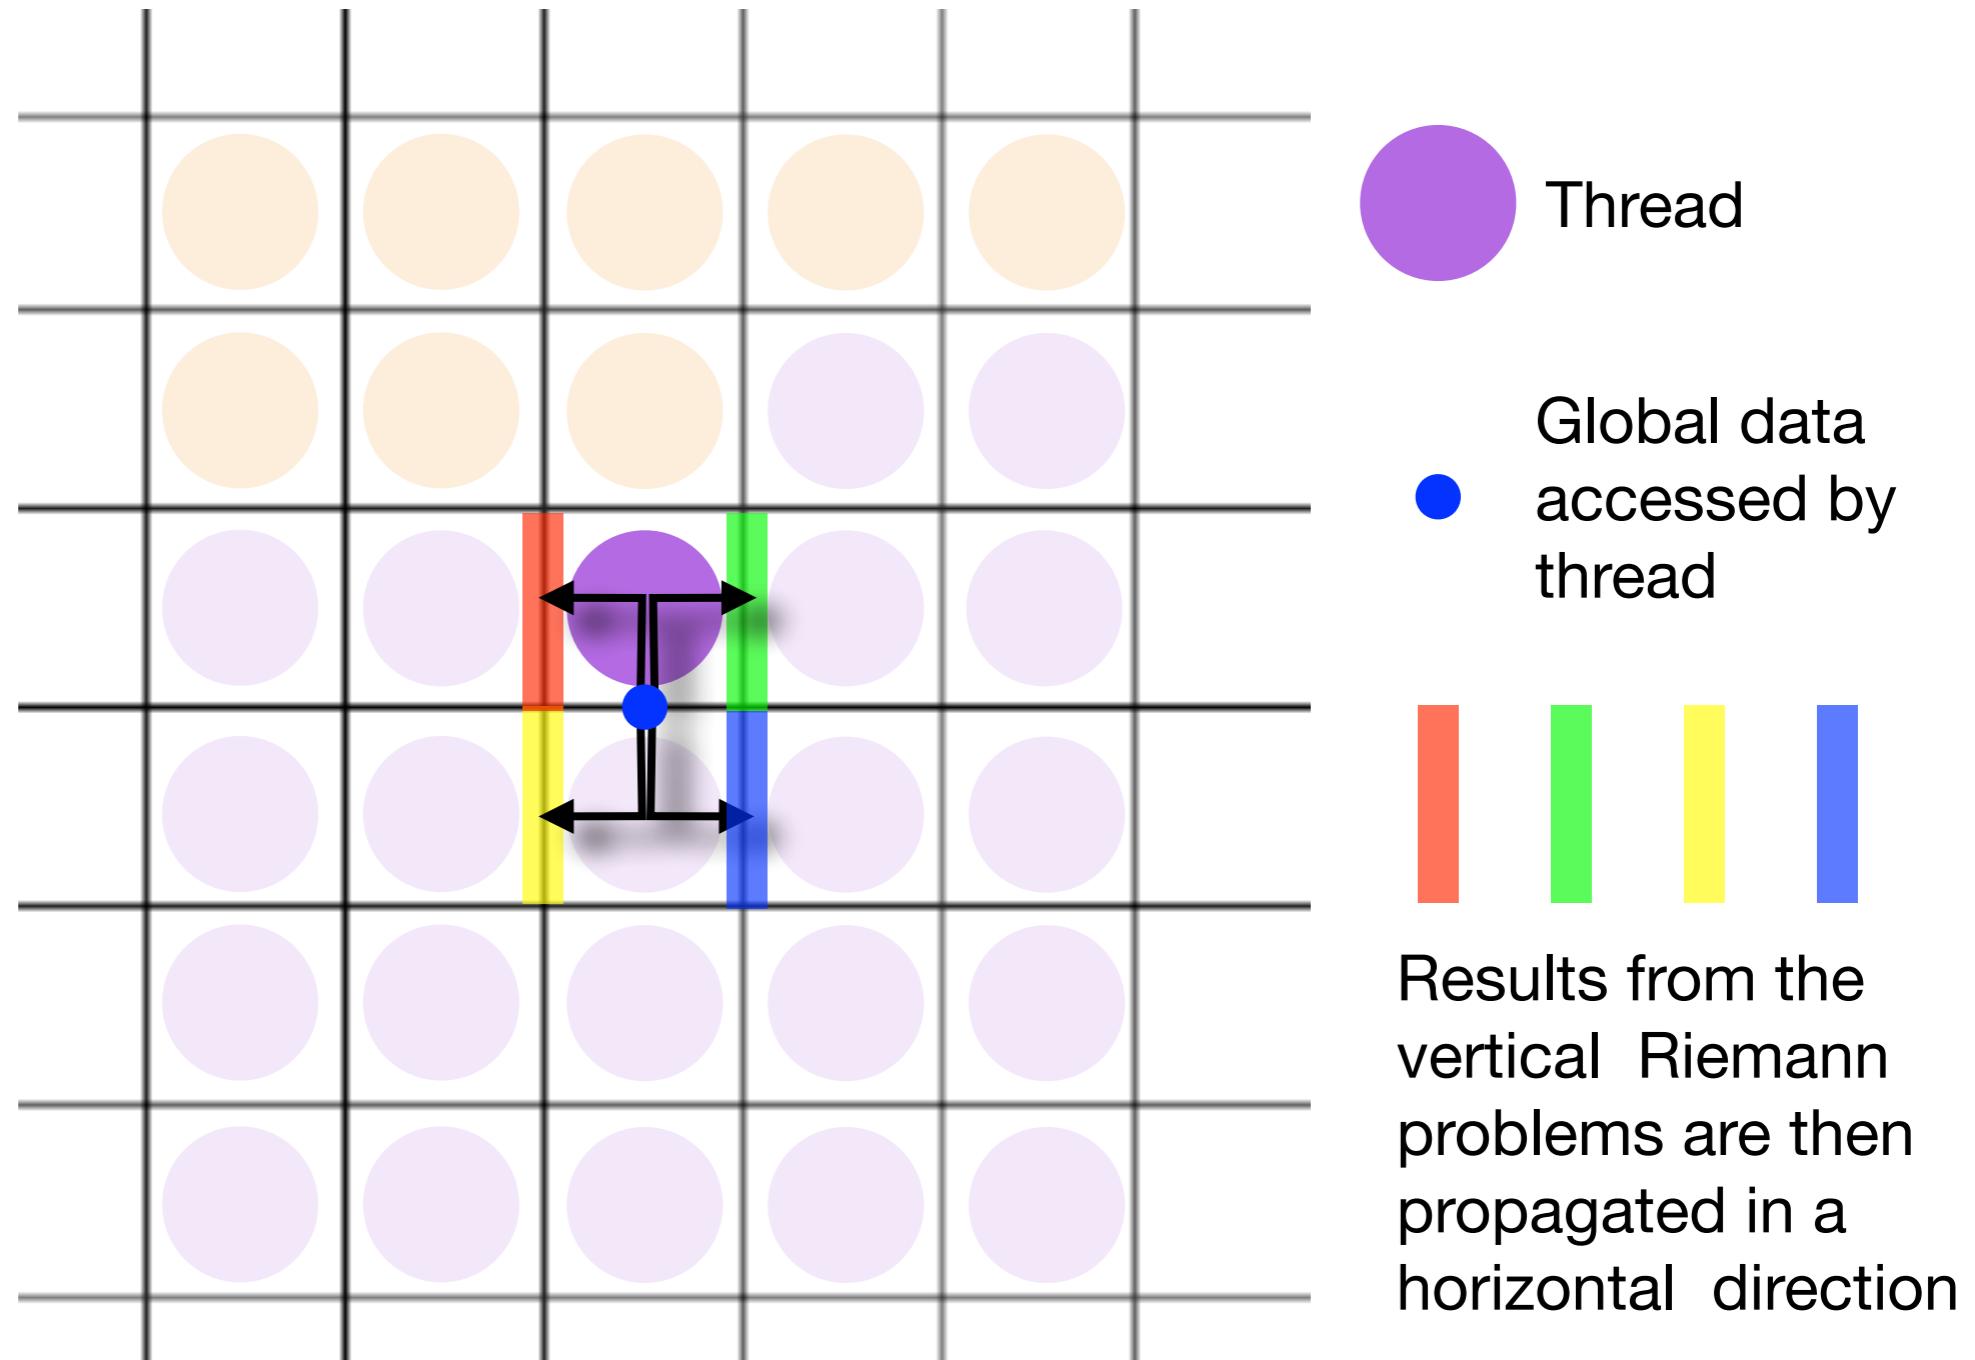
- Each transverse solve stores data in the same global memory space
- To avoid data collisions with other threads writing to the same global memory, four passes over all the global data are required, one for each “color”
- Sync threads between each pass



Fluxes are computed by solving Riemann problems

Transverse Riemann problems

- Each transverse solve stores data in the same global memory space
- Four more passes over all the global data are required

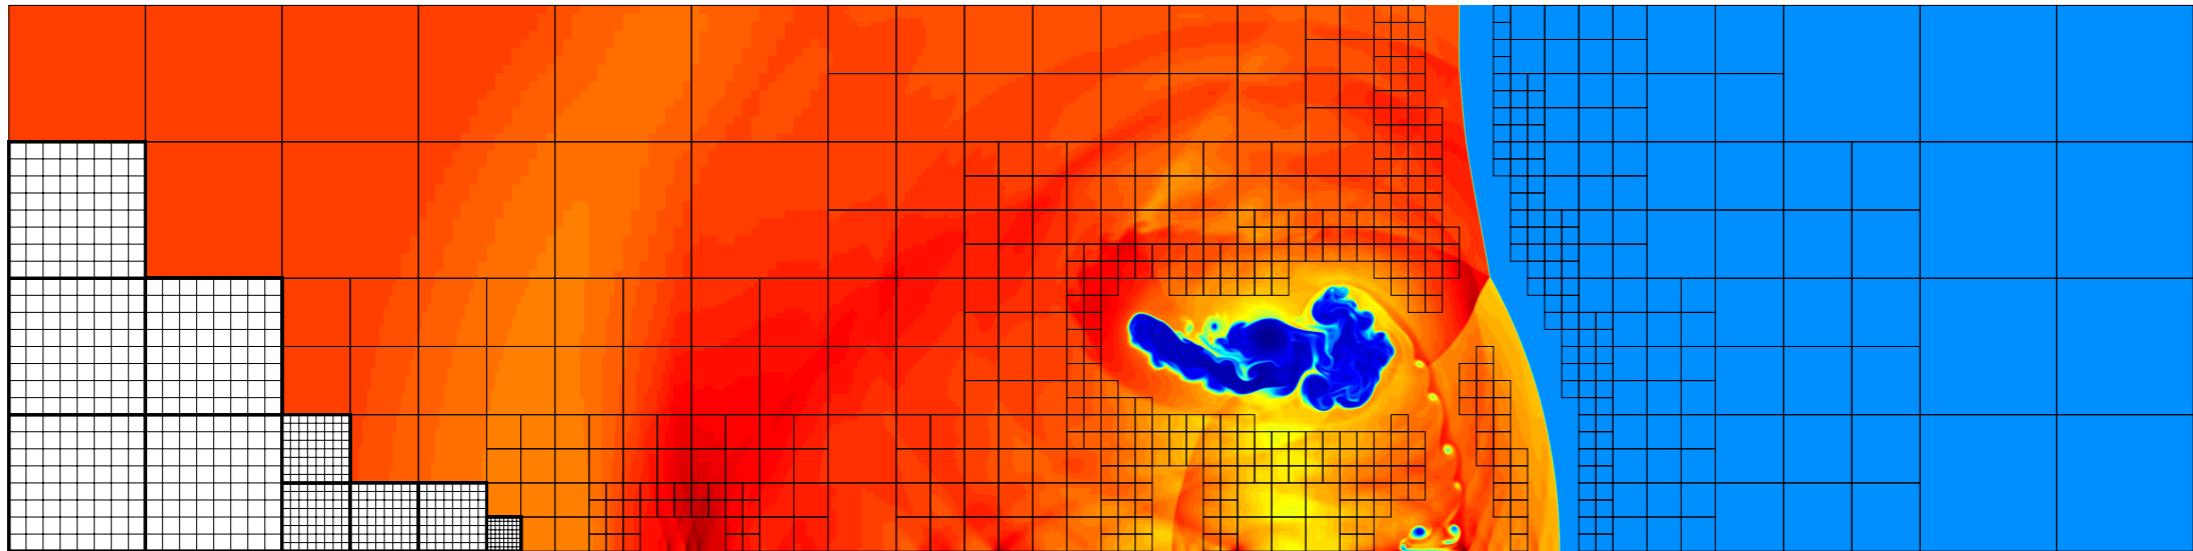


Unsplit wave propagation

Fully unsplit wave propagation algorithm is implemented in a single CUDA kernel.

- While more expensive than the dimensionally split version, the unsplit algorithm may be more suited to AMR.
- The cost in CUDA is that parts of the code that can be done together are now split to avoid race conditions. *Maybe we can improve on this by using more global memory?*
- Our GPU configuration does not require any synchronization between thread blocks
- Since all patches are the same size, they can be processed in large batches ($O(1000)$ per batch)
- All AMR tasks including filling ghost cells is done on the CPU
- Conservation requires extra memory copy from device.

Shock-bubble problem



- Euler equations : Four field variables per finite volume cell
- Riemann solvers written in CUDA and passed in as CUDA pointers
- 32x32 patch sizes seem optimal
- Ran on 4 node (4 cores per node) cluster with 2 GeForce Titan X (2015) per node
- CPU and GPU results agree to machine precision

Related work :

H. G. Ohannessian, G. Turkiyyah, A. J. Ahmadia, and D. I. Ketcheson, CUDACLAW: A high-performance programmable GPU framework for the solution of hyperbolic PDEs, arXiv, 1805.08846 (2018).

X. Qin, R. J. LeVeque, M. Motley, “Accelerating wave-propagation algorithms on adaptive mesh with the graphics processing unit (GPU)”, 2018.

EuroHack 2018 - Lugano, Switzerland



Scott Aiton (BSU), Andreas Jocksch (CSCS), Xinsheng Qin (Univ. of Washington), D. Calhoun (BSU), Melody Shih (NYU)

Sponsored by : NVIDIA + Swiss National Computing Center